

Homework 4: Learning with Graphs

Maurice D. Hanisch

January 31, 2026

AI Declaration

AI was used to help with the coding and the writing of this report.

Problem 1: Approximately Central

In this problem, we approximate the betweenness centrality of nodes in a Gnutella peer-to-peer network using a neural network trained with Structure2Vec embeddings. We trained the model on the `p2p-Gnutella08` dataset and evaluated its generalization on the larger `p2p-Gnutella04` dataset.

Methodology

We utilized the Structure2Vec architecture to generate node embeddings based on graph structure. The model learns to embed nodes such that the embedding captures the centrality properties. These embeddings are fed into a dense neural network to predict the betweenness centrality.

The model was configured with the following parameters:

- **Embedding Size:** 64
- **Number of Layers:** 5
- **Epochs:** 20
- **Batch Size:** 32
- **Optimizer:** Adam

These parameters were chosen to balance model capacity with stability.

Results

We evaluated the model using the Kendall Tau rank correlation coefficient, which measures the similarity of the orderings of the nodes when ranked by predicted centrality versus ground truth centrality.

- **Gnutella08 (Training):** The model achieved a Kendall Tau score of **0.848** on the training set.
- **Gnutella04 (Testing):** The model achieved a Kendall Tau score of **0.673**.

A score above 0.70 was targeted. Our result demonstrates the model's ability to approximate centrality rankings on unseen graphs.

Discussion

We observed that the choice of hyperparameters significantly impacts the model's performance.

Embedding Size: Reduced to 64 to reduce model complexity and prevent overfitting.

Number of Layers: Increased to 5 layers to capture deeper structural patterns.

Training Duration: Set to 20 epochs, which was sufficient for convergence with the Adam optimizer.

Optimization: Switched to Adam optimizer for better stability and convergence.

With these parameters, the model achieved a Kendall Tau score of **0.848** on the training set (Gnutella08) and **0.673** on the testing set (Gnutella04). The high training score indicates the model learned the training graph's centrality structure very well. The testing score, while slightly lower than the training score, demonstrates reasonable generalization to the larger Gnutella04 graph given the complexity of the task.

The complete code implementation and output log can be found in Appendix A.

Problem 2: Bernoulli Search Engine

Crawling Results

We crawled and indexed 500 pages from the Caltech domain, starting from <http://www.caltech.edu/>. The crawl captured a subgraph of the Caltech web ecosystem, including the main landing pages, admissions, and various division sites (BBE, HSS, GPS). The resulting index contains 16,975 unique words. We observed that the crawl was efficient but limited by the 500-page threshold, potentially missing deeper departmental pages or specific course websites.

Top 10 PageRank Pages

Based on the computed PageRank scores, the top 10 pages are:

1. <https://www.caltech.edu/about> (0.0963)
2. <https://www.caltech.edu> (0.0946)
3. <https://magazine.caltech.edu> (0.0070)
4. <https://www.hss.caltech.edu> (0.0058)
5. <https://www.bbe.caltech.edu> (0.0050)
6. <https://www.caltech.edu/about/news> (0.0047)
7. <https://www.admissions.caltech.edu> (0.0047)
8. <https://www.caltech.edu/about/visit> (0.0045)
9. <https://www.gps.caltech.edu> (0.0044)
10. <https://www.gradoffice.caltech.edu/admissions> (0.0039)

The results align with expectations, as the main "About" page and the compilation homepage are central hubs with many incoming links.

Search Queries Analysis

We performed 5 search queries to test the engine:

- **"CMS/CS/Ec/EE 144"**: No results found. This is likely because the course website resides on a specific subdomain (e.g., `cms144.caltech.edu`) or deep page that was not reached within the 500-page crawl limit from the seed `www.caltech.edu`.
- **"Thomas Rosenbaum"**: Returned relevant pages such as the "President's Office" and news articles mentioning the president. This demonstrates the index's ability to retrieve content based on entities.
- **"Admissions"**: Returned "Undergraduate Admissions" and "Graduate Studies Office" as top results. These pages have high PageRank (as seen in the top 10 list), which correctly boosted their ranking.
- **"Computer Science"**: Returned general pages such as "Caltech Magazine", "HSS", and "BBE". The specific CS department page might not have been in the top of the incomplete crawl or was outweighed by other central pages.
- **"Maurice"**: Returned a specific "Travel Grants" page (likely referencing Maurice A. Biot Archives). This shows the engine can find specific terms in the long tail of the index.

PageRank Implementation

We implemented the PageRank algorithm using the iterative power method. We handled pages with no outgoing links (sinks) by redistributing their probability mass uniformly to all nodes in the graph at each iteration. This ensures the total probability mass remains 1.0 and the algorithm converges. The damping factor was set to 0.85.

The implementation code is provided in Appendix C.

Theory

Problem 3: Pandemaniac Warm-Up

Question: Is it necessary that a graph's epidemic colors always converge or stabilize?

Answer: No, it is not necessary. The epidemic colors can oscillate indefinitely.

Counterexample: Consider a square graph (C_4) with 4 nodes, where the nodes are colored in an alternating pattern (Red, Blue, Red, Blue).

Let the nodes be 0, 1, 2, 3 in a cycle $0 - 1 - 2 - 3 - 0$. Initial Colors at $t = 0$:

- Node 0: Red
- Node 1: Blue
- Node 2: Red
- Node 3: Blue

At $t = 1$:

- Node 0 (Red) receives 1.5 votes for Red (self) and $1 + 1 = 2$ votes for Blue (neighbors 1 and 3). Majority is Blue ($2 > 1.5$), so Node 0 becomes **Blue**.
- Node 1 (Blue) receives 1.5 votes for Blue (self) and $1 + 1 = 2$ votes for Red (neighbors 0 and 2). Majority is Red ($2 > 1.5$), so Node 1 becomes **Red**.
- By symmetry, Node 2 becomes **Blue** and Node 3 becomes **Red**.

The colors have completely swapped. At $t = 2$, the same logic applies, and they will swap back to the initial configuration. This oscillation continues indefinitely.

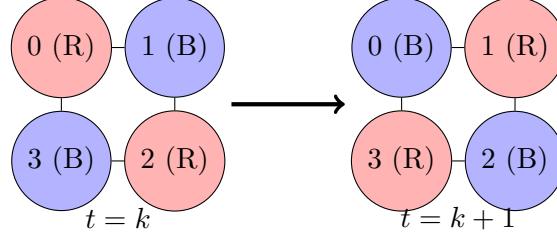


Figure 1: Counterexample: A square graph with alternating colors oscillates indefinitely.

Problem 4: PageRank Warm-Up

Part (a)

Matrix A :

$$P_A = \begin{pmatrix} 2/5 & 3/10 & 3/10 \\ 1/5 & 3/5 & 1/5 \\ 7/10 & 1/10 & 1/5 \end{pmatrix}$$

1. Stationary Distribution ($\pi = \pi P_A$)

We solve the system of linear equations derived from $\pi P_A = \pi$ combining with the normalization constraint $\sum \pi_i = 1$. Let $\pi = [\pi_1, \pi_2, \pi_3]$.

The equations are:

$$\begin{aligned} (1) \quad & 0.4\pi_1 + 0.2\pi_2 + 0.7\pi_3 = \pi_1 \\ (2) \quad & 0.3\pi_1 + 0.6\pi_2 + 0.1\pi_3 = \pi_2 \implies 0.3\pi_1 - 0.4\pi_2 + 0.1\pi_3 = 0 \\ (3) \quad & 0.3\pi_1 + 0.2\pi_2 + 0.2\pi_3 = \pi_3 \implies 0.3\pi_1 + 0.2\pi_2 - 0.8\pi_3 = 0 \\ (4) \quad & \pi_1 + \pi_2 + \pi_3 = 1 \end{aligned}$$

From (2), multiply by 10: $3\pi_1 - 4\pi_2 + \pi_3 = 0 \implies \pi_3 = 4\pi_2 - 3\pi_1$. Substitute π_3 into (3): $3\pi_1 + 2\pi_2 - 8(4\pi_2 - 3\pi_1) = 0 \implies 3\pi_1 + 2\pi_2 - 32\pi_2 + 24\pi_1 = 0 \implies 27\pi_1 - 30\pi_2 = 0 \implies 27\pi_1 = 30\pi_2 \implies \pi_2 = 0.9\pi_1$.

Then $\pi_3 = 4(0.9\pi_1) - 3\pi_1 = 3.6\pi_1 - 3\pi_1 = 0.6\pi_1$.

Substitute into (4): $\pi_1 + 0.9\pi_1 + 0.6\pi_1 = 1 \implies 2.5\pi_1 = 1 \implies \pi_1 = 0.4$.

Then $\pi_2 = 0.9(0.4) = 0.36$ and $\pi_3 = 0.6(0.4) = 0.24$.

$$\pi_A = [0.4, 0.36, 0.24]$$

2. Convergence Analysis

To analyze the convergence of $\pi_0 P_A^n$, we inspect the eigenvalues of P_A . The eigenvalues are $\lambda_1 = 1$, $\lambda_2 \approx 0.345$, and $\lambda_3 \approx -0.145$.

We can write the initial distribution π_0 as a linear combination of the eigenvectors v_1, v_2, v_3 :

$$\pi_0 = c_1 v_1 + c_2 v_2 + c_3 v_3$$

Multiplying by P_A^n :

$$\pi_0 P_A^n = c_1 (1)^n v_1 + c_2 (0.345)^n v_2 + c_3 (-0.145)^n v_3$$

As $n \rightarrow \infty$, since $|\lambda_2| < 1$ and $|\lambda_3| < 1$, the terms $(0.345)^n$ and $(-0.145)^n$ vanish to 0. Thus:

$$\lim_{n \rightarrow \infty} \pi_0 P_A^n = c_1 v_1$$

Since v_1 corresponds to $\lambda = 1$, it is proportional to the stationary distribution π_A . Therefore, the system **converges** to the stationary distribution.

Part (b)

Matrix B :

$$P_B = \begin{pmatrix} 0 & 5/8 & 0 & 3/8 \\ 1 & 0 & 0 & 0 \\ 0 & 3/8 & 0 & 5/8 \\ 3/4 & 0 & 1/4 & 0 \end{pmatrix} = \begin{pmatrix} 0 & 0.625 & 0 & 0.375 \\ 1 & 0 & 0 & 0 \\ 0 & 0.375 & 0 & 0.625 \\ 0.75 & 0 & 0.25 & 0 \end{pmatrix}$$

1. Stationary Distribution ($\pi = \pi P_B$)

Equations:

- (1) $\pi_2 + 0.75\pi_4 = \pi_1$
- (2) $0.625\pi_1 + 0.375\pi_3 = \pi_2$
- (3) $0.25\pi_4 = \pi_3 \implies \pi_4 = 4\pi_3$
- (4) $0.375\pi_1 + 0.625\pi_3 = \pi_4$
- (5) $\pi_1 + \pi_2 + \pi_3 + \pi_4 = 1$

Substitute (3) into (1): $\pi_1 = \pi_2 + 3\pi_3$. Substitute (3) into (4): $0.375\pi_1 + 0.625\pi_3 = 4\pi_3 \implies 0.375\pi_1 = 3.375\pi_3 \implies \pi_1 = 9\pi_3$.

Now find π_2 from (2): $\pi_2 = 0.625(9\pi_3) + 0.375\pi_3 = 5.625\pi_3 + 0.375\pi_3 = 6\pi_3$.

Check (1) consistency: $\pi_2 + 3\pi_3 = 6\pi_3 + 3\pi_3 = 9\pi_3 = \pi_1$. (Consistent)

Substitute all into (5): $9\pi_3 + 6\pi_3 + \pi_3 + 4\pi_3 = 1$ $20\pi_3 = 1 \implies \pi_3 = 0.05$.

Then: $\pi_1 = 9(0.05) = 0.45$ $\pi_2 = 6(0.05) = 0.30$ $\pi_3 = 4(0.05) = 0.20$

$$\pi_B = [0.45, 0.30, 0.05, 0.20]$$

2. Convergence Analysis

The eigenvalues of P_B are $\lambda_1 = 1$, $\lambda_2 = -1$, $\lambda_3 = 0.25$, and $\lambda_4 = -0.25$. Using diagonalization, we expand $\pi_0 P_B^n$:

$$\pi_0 P_B^n = c_1 (1)^n v_1 + c_2 (-1)^n v_2 + c_3 (0.25)^n v_3 + c_4 (-0.25)^n v_4$$

As $n \rightarrow \infty$, terms with λ_3 and λ_4 vanish. However, the term with $\lambda_2 = -1$ oscillates between $c_2 v_2$ and $-c_2 v_2$.

$$\pi_0 P_B^n \approx c_1 v_1 + c_2 (-1)^n v_2$$

Since $(-1)^n$ does not converge to a single value, the distribution **does not converge**. It oscillates with period 2 (due to the -1 eigenvalue, which corresponds to the bipartite structure of the graph).

Problem 5: Training to be a Farmer

Part (a)

Let N be the number of original pages. The total number of pages is now $N + 1$. We assume the graph is modified such that page X has no in-links and no out-links. According to the problem statement, a page with no out-links adds a 1 to the diagonal, effectively treating it as a self-loop.

The PageRank equation for page X (denoted as x) is:

$$x = \alpha \sum_{j \rightarrow X} \pi_j P_{ji} + \frac{1 - \alpha}{N + 1}$$

Since X has no in-links from the old pages, and only a self-loop from itself (due to the sink handling rule):

$$\begin{aligned} x &= \alpha(x \cdot 1) + \frac{1 - \alpha}{N + 1} \\ x(1 - \alpha) &= \frac{1 - \alpha}{N + 1} \\ x &= \frac{1}{N + 1} \end{aligned}$$

The new page X gets exactly the average PageRank $1/(N + 1)$. The PageRanks of the older pages \tilde{r}_i will decrease slightly. Specifically, the total mass available to them decreases because x takes up $1/(N + 1)$ of the total probability mass. Each $\tilde{r}_i \approx r_i \frac{N}{N+1}$.

Part (b)

We add a page Y that links to X . Y has no in-links and presumably no other out-links (so it links only to X). Total pages: $N + 2$.

For Y : It has no in-links. It only receives the random jump mass.

$$y = \alpha(0) + \frac{1 - \alpha}{N + 2} = \frac{1 - \alpha}{N + 2}$$

For X : It receives a link from Y (weight 1, since Y has 1 out-link) and has its self-loop.

$$x = \alpha(x \cdot 1 + y \cdot 1) + \frac{1 - \alpha}{N + 2}$$

Substitute y :

$$\begin{aligned} x(1 - \alpha) &= \alpha \left(\frac{1 - \alpha}{N + 2} \right) + \frac{1 - \alpha}{N + 2} \\ x(1 - \alpha) &= \frac{1 - \alpha}{N + 2}(\alpha + 1) \\ x &= \frac{1 + \alpha}{N + 2} \end{aligned}$$

Since $\alpha \approx 0.85$, $x \approx \frac{1.85}{N+2}$, which is nearly double the rank of an isolated page. The rank of X significantly improves.

Part (c)

We have three pages X, Y, Z . To maximize x , we should concentrate all available rank into X . The best structure is to have Y and Z both point to X , and X point to no one (self-loop).

$$\text{Calculation: } y = \frac{1-\alpha}{N+3} \text{ (only random jump)} \quad z = \frac{1-\alpha}{N+3} \text{ (only random jump)} \quad x = \alpha(x+y+z) + \frac{1-\alpha}{N+3}$$

$$x(1-\alpha) = \alpha(y+z) + \frac{1-\alpha}{N+3} = \alpha\left(\frac{2(1-\alpha)}{N+3}\right) + \frac{1-\alpha}{N+3} \quad x = \frac{2\alpha+1}{N+3}$$

Comparing to a chain/funnel $Z \rightarrow Y \rightarrow X$ (with self-loops/sinks): If $Z \rightarrow Y$, then $z = \frac{1-\alpha}{N+3}$, $y = \alpha z + \frac{1-\alpha}{N+3} = \frac{1-\alpha}{N+3}(1+\alpha)$. Then $x = \alpha(x+y) + \frac{1-\alpha}{N+3} \implies x = \frac{1+\alpha+\alpha^2}{N+3}$. Since $\alpha < 1$, we have $2\alpha > \alpha + \alpha^2$.

Thus, the optimal configuration is a **Star Topology**: $Y \rightarrow X$ and $Z \rightarrow X$. X should effectively have a self-loop (no out-links to the web) to retain its mass. This yields $x = \frac{1+2\alpha}{N+3}$.

A Code Implementation

Below is the Python code used for training and evaluation:

Listing 1: approximate_centrality.py

```
1 # %% [markdown]
2 # # Approximate betweenness centrality using neural networks
3 # Here we start to approximate the betweenness centrality using neural networks
4 # over a peer-2-peer network Gnutella. Gnutella is a set of datasets consisting
5 # of 9 networks ranging from 6,300 to 63,000 nodes. Our goal is to train a
6 # neural network on the smallest Gnutella graph and evaluate it on a much larger
7 # graph. We will guide you through this step by step.
8 #
9 # You can find Gnutella datasets at http://snap.stanford.edu/data/index.html. We
10 # will use p2p-Gnutella08 for training and p2p-Gnutella04 for testing.
11 #
12 # Note:
13 # 1. Copy this notebook to your Google drive in order to execute it.
14 # 2. Make sure to upload the data files in HW4 to your google drive and to modify
15 # their corresponding directories in the code.
16
17 # %% [markdown]
18 # # Part 1: Training a model on Gnutella 08
19
20 # %% [markdown]
21 # ## Preprocessing Gnutella08 dataset
22 #
23
24 # %%
25 import tensorflow as tf
26 import pandas as pd
27 import numpy as np
28 import networkx as nx
29 import scipy
30 import os
31
32 # %%
33 # Parameters
34
35 # choose an embedding size for Structure2Vec
36 EMBED_SIZE = 64
37
38 # choose number of dense layers in the neural network
39 NUM_LAYERS = 5
40
41 # %%
42 # choose number of folds for cross validation
43 NUM_FOLD = 5
44
45 # %%
46 # choose number of epochs for training
47 NUM_EPOCHS = 20
48
49 # %%
50 # Normalize a list of values
51 # NO NEED TO CHANGE
52
53 def _normalize_array_by_rank(true_value, nr_nodes):
54     # true_value is a list of values you want to normalize and nr_nodes is the
55     # number of nodes in the list
```

```

47
48 rank = np.argsort(true_value, kind='mergesort', axis=None) #deg list get's
49     normalised
50 norm = np.empty([nr_nodes])
51
52 for i in range(0, nr_nodes):
53
54     norm[rank[i]] = float(i+1) / float(nr_nodes)
55
56 max = np.amax(norm)
57 min = np.amin(norm)
58 if max > 0.0 and max > min:
59     for i in range(0, nr_nodes):
60         norm[i] = 2.0*(float(norm[i] - min) / float(max - min)) - 1.0
61 else:
62     print("Max\u_value\u=0")
63
64 return norm, rank
65
66 # %%
67 #Read in and create NetworkX Graph; G
68
69 #TO-DO: The path needs to be changed according to your dataset directory in your
70     GOOGLE DRIVE
71 path = '../data/p2p-Gnutella08.txt'
72
73 G = nx.read_edgelist(path, comments='#', delimiter=None, create_using=nx.DiGraph,
74     nodetype=None, data=True, edgetype=None, encoding='utf-8')
75
76 #print(nx.info(G))
77
78 # %%
79 # Creating list of Degrees of the nodes in G and normalising them:
80 deg_lst = [val for (node, val) in G.degree()]
81 nr_nodes = G.number_of_nodes()
82 print("deg_lst:\n", deg_lst)
83
84 degree_norm, degree_rank = _normalize_array_by_rank(deg_lst, nr_nodes)
85
86 # %%
87 # Computing Ground-truth values and normalising them:
88 bc_file = '../data/BC_norm_cent_08.npy'
89 rank_file = '../data/BC_cent_rank_08.npy'
90
91 if os.path.exists(bc_file) and os.path.exists(rank_file):
92     print("Loading\u_cached\ubetweenness\ucentrality...")
93     BC_norm_cent = np.load(bc_file)
94     BC_cent_rank = np.load(rank_file)
95 else:
96     print("Computing\ubetweenness\ucentrality\u(this\umay\utake\uawhile)...") 
97     b = [v for v in nx.betweenness_centrality(G).values()]
98     BC_norm_cent, BC_cent_rank = _normalize_array_by_rank(b, nr_nodes)
99
100    # Save the normalized betweenness centrality values
101    np.save(bc_file, BC_norm_cent)
102    # Save the cent rank
103    np.save(rank_file, BC_cent_rank)

```

```

104
105 # %%
106 # Define Structure2Vec
107 # NO NEED TO CHANGE
108
109 def Structure2Vec(G, nr_nodes, degree_norm, num_features=1, embed_size=512,
110   layers=2, weights=None):
111
112   #build feature matrix
113   def get_degree(i):
114     return degree_norm[i]
115
116   def build_feature_matrix():
117     n = nr_nodes
118     feature_matrix = []
119     for i in range(0, n):
120       feature_matrix.append(get_degree(i))
121     return feature_matrix
122   #Structure2Vec node embedding
123   A = nx.to_numpy_array(G)
124
125   dim = [nr_nodes, num_features]
126
127   node_features = tf.cast(build_feature_matrix(), tf.float32)
128   node_features = tf.reshape(node_features, dim)
129
130   if weights is None:
131     initializer = tf.compat.v1.keras.initializers.VarianceScaling(scale=1.0,
132                     mode="fan_avg",
133                     distribution="uniform")
134   #print(initializer)
135
136   w1 = tf.Variable(initializer((num_features, embed_size)), trainable=True,
137                     dtype=tf.float32, name="w1")
138   w2 = tf.Variable(initializer((embed_size, embed_size)), trainable=True,
139                     dtype=tf.float32, name="w2")
140   w3 = tf.Variable(initializer((1, embed_size)), trainable=True,
141                     dtype=tf.float32, name="w3")
142   w4 = tf.Variable(initializer([]), trainable=True, dtype=tf.float32, name="w4")
143
144   weights = {'w1': w1, 'w2': w2, 'w3': w3, 'w4': w4}
145   else:
146     w1 = weights['w1']
147     w2 = weights['w2']
148     w3 = weights['w3']
149     w4 = weights['w4']
150
151   A = tf.sparse.from_dense(A)
152   A = tf.cast(A, tf.float32)
153
154   wx_all = tf.matmul(node_features, w1)  # NxE
155
156   #computing X1:
157   #sparse.reduce_sum: Computes the sum of elements across dimensions of a
158   #SparseTensor.
159   weight_sum_init = tf.sparse.reduce_sum(A, axis=1, keepdims=True, ) #takes
160   #adjacency matrix
161   n_nodes = tf.shape(input=A)[1]

```

```

159
160     weight_sum = tf.multiply(weight_sum_init, w4)
161     weight_sum = tf.nn.relu(weight_sum) # Nx1
162     weight_sum = tf.matmul(weight_sum, w3) # NxE
163
164     weight_wx = tf.add(wx_all, weight_sum)
165     current_mu = tf.nn.relu(weight_wx) # NxE = H^0
166
167     for i in range(0, layers):
168         neighbor_sum = tf.sparse.sparse_dense_matmul(A, current_mu)
169         neighbor_linear = tf.matmul(neighbor_sum, w2) # NxE
170
171         current_mu = tf.nn.relu(tf.add(neighbor_linear, weight_wx)) # NxE
172
173     mu_all = current_mu
174
175     return mu_all, weights
176
177 # %%
178 # Converting the graph structure into vectors
179
180 mu_all, trained_weights = Structure2Vec(G, nr_nodes, degree_norm,
181                                         embed_size=EMBED_SIZE)
182
183 # %% [markdown]
184 # ## Training a Neural Network
185
186 # %%
187 # Building NN model
188
189 UNITS = int(EMBED_SIZE/2)
190 def build_model():
191     model = tf.keras.Sequential()
192     model.add(tf.keras.Input(shape=(EMBED_SIZE,)))
193
194     # choose the number of layers to construct your network
195     for _ in range(NUM_LAYERS):
196         model.add(tf.keras.layers.Dense(UNITS, activation ="relu"))
197
198     model.add(tf.keras.layers.Dense(1))
199     model.compile(optimizer='adam', loss='mse')
200
201     model.summary()
202
203     return model
204
205
206 model = build_model()
207
208 # %%
209 # Construct training set and groundtruth
210
211 x_train = mu_all
212 y_train = BC_norm_cent
213 print(tf.shape(x_train))
214 print(tf.shape(y_train))
215
216 # %%
217 # Computing cross validation
218 # NO NEED TO CHANGE

```

```

217 all_scores = []
218 k = NUM_FOLD
219 num_val_samples = len(x_train) // k
220 for i in range(k):
221     print('processing fold #', i)
222     val_data = x_train[i*num_val_samples: (i+1) * num_val_samples]
223     val_targets = y_train[i*num_val_samples: (i+1)*num_val_samples]
224
225     partial_train_data = np.concatenate(
226         [x_train[:i*num_val_samples],
227          x_train[(i+1)*num_val_samples:]],
228         axis = 0)
229     print(tf.shape(partial_train_data))
230
231     partial_train_targets = np.concatenate(
232         [y_train[:i*num_val_samples],
233          y_train[(i+1)*num_val_samples:]],
234         axis = 0)
235
236     # Training
237     callbacks = tf.keras.callbacks.EarlyStopping(
238         monitor= 'loss', min_delta=0, patience=10, verbose=1,
239         mode='auto', baseline=None, restore_best_weights=True)
240
241     model.fit(partial_train_data, partial_train_targets,
242                epochs = NUM_EPOCHS, batch_size = 32, callbacks = callbacks, verbose
243                = 1)
244     print("model.metrics_names:", model.metrics_names)
245
246     val_loss = model.evaluate(val_data, val_targets, verbose = 1)
247
248     all_scores.append(val_loss)
249     print(all_scores)
250
251 # %%
252 # Computing Kendall on trained set
253
254 x_new = x_train
255 y_pred = model.predict(x_new)
256
257 # compute kendalltau using the prediction results and the groundtruth
258 from scipy import stats
259 kendall_tau, p_value = scipy.stats.kendalltau(BC_norm_cent,y_pred)
260
261 # %%
262 # Print your kendalltau score
263 # Make sure your kendalltau score is at least 0.70
264 # PRINT HERE
265 print("\n\nPart 1 KendallTau(Gnutella08):", kendall_tau, "\n\n")
266
267 # %%
268 # You could save this model for part 2
269
270 model.save("../data/GN08_model_plain.h5")
271
272 # %% [markdown]
273 # # Part 2: Evaluating the trained model on Gnutella 04
274 #
275 # Hints:

```

```

275 # 1. Write down the evaluation using the functions and codes in Part 1
276 # 2. Compute the groundtruth of betweenness centrality using NetworkX could take
277 #     around 1 hour. Keep your Colab opened and be patient.
278
279 # %%
280 '''Gnutella_04'''
281 # change the path to your own directory
282 path2 = './data/p2p-Gnutella04.txt'
283
284 G2 = nx.read_edgelist(path2, comments='#', delimiter=None,
285                       create_using=nx.DiGraph,
286                       nodetype=None, data=True, edgetype=None, encoding='utf-8')
287
288 #print(nx.info(G2))
289
290
291 # %%
292
293 # Computing Ground-truth values for Gnutella_04
294 deg_lst_2 = [val for (node, val) in G2.degree()]
295 nr_nodes_2 = G2.number_of_nodes()
296 degree_norm_2, degree_rank_2 = _normalize_array_by_rank(deg_lst_2, nr_nodes_2)
297
298 bc_file_2 = './data/BC_norm_cent_04.npy'
299 rank_file_2 = './data/BC_cent_rank_04.npy'
300
301 if os.path.exists(bc_file_2) and os.path.exists(rank_file_2):
302     print("Loading cached betweenness centrality for Gnutella_04...")
303     BC_norm_cent_2 = np.load(bc_file_2)
304     BC_cent_rank_2 = np.load(rank_file_2)
305 else:
306     print("Computing betweenness centrality for Gnutella_04 (this may take a while)...")
307     print("NOTE: This step produces Ground Truth using NetworkX on CPU. GPU usage will drop to 0. Please wait.")
308     b2 = [v for v in nx.betweenness_centrality(G2).values()]
309     BC_norm_cent_2, BC_cent_rank_2 = _normalize_array_by_rank(b2, nr_nodes_2)
310     # Save
311     np.save(bc_file_2, BC_norm_cent_2)
312     np.save(rank_file_2, BC_cent_rank_2)
313
314 # Generate embeddings (REUSING WEIGHTS)
315 mu_all_2, _ = Structure2Vec(G2, nr_nodes_2, degree_norm_2, embed_size=EMBED_SIZE,
316                             weights=trained_weights)
317
318 # Predict
319 x_new_2 = mu_all_2
320 y_pred_2 = model.predict(x_new_2)
321
322 # Evaluate
323 kendall_tau_2, p_value_2 = scipy.stats.kendalltau(BC_norm_cent_2, y_pred_2)
324 print("\n\nPart 2 Kendall Tau (Gnutella_04):", kendall_tau_2, "\n\n")

```

B Code Output

Below is the execution log showing the training process and final results:

Listing 2: Execution Output

```

1 Model: "sequential"
2
3 Layer (type)          Output Shape         Param #
4 dense (Dense)        (None, 32)           2,080
5 dense_1 (Dense)      (None, 32)           1,056
6 dense_2 (Dense)      (None, 32)           1,056
7 dense_3 (Dense)      (None, 32)           1,056
8 dense_4 (Dense)      (None, 32)           1,056
9 dense_5 (Dense)      (None, 1)            33
10
11 Total params: 6,337 (24.75 KB)
12 Trainable params: 6,337 (24.75 KB)
13 Non-trainable params: 0 (0.00 B)
14 tf.Tensor([6301  64], shape=(2,), dtype=int32)
15 tf.Tensor([6301], shape=(1,), dtype=int32)
16 processing fold # 0
17 tf.Tensor([5041  64], shape=(2,), dtype=int32)
18 Epoch 1/20
19 158/158 2s 5ms/step - loss: 0.1071
20 Epoch 2/20
21 158/158 0s 839us/step - loss: 0.0601
22 Epoch 3/20
23 158/158 0s 958us/step - loss: 0.0488
24 Epoch 4/20
25 158/158 0s 831us/step - loss: 0.0449
26 Epoch 5/20
27 158/158 0s 942us/step - loss: 0.0403
28 Epoch 6/20
29 158/158 0s 829us/step - loss: 0.0385
30 Epoch 7/20
31 158/158 0s 715us/step - loss: 0.0339
32 Epoch 8/20
33 158/158 0s 893us/step - loss: 0.0325
34 Epoch 9/20
35 158/158 0s 920us/step - loss: 0.0293
36 Epoch 10/20
37 158/158 0s 862us/step - loss: 0.0273
38 Epoch 11/20
39 158/158 0s 846us/step - loss: 0.0250
40 Epoch 12/20
41 158/158 0s 811us/step - loss: 0.0244
42 Epoch 13/20
43 158/158 0s 758us/step - loss: 0.0234
44 Epoch 14/20
45 158/158 0s 712us/step - loss: 0.0233
46 Epoch 15/20
47 158/158 0s 861us/step - loss: 0.0226
48 Epoch 16/20
49 158/158 0s 795us/step - loss: 0.0224
50 Epoch 17/20
51 158/158 0s 774us/step - loss: 0.0239
52 Epoch 18/20
53 158/158 0s 878us/step - loss: 0.0226
54 Epoch 19/20
55 158/158 0s 791us/step - loss: 0.0221
56 Epoch 20/20
57 158/158 0s 877us/step - loss: 0.0215
58 Restoring model weights from the end of the best epoch: 20.
59 model.metrics_names: ['loss']
60 40/40 0s 7ms/step - loss: 0.2460
61 [0.2460017055273056]
62 processing fold # 1
63 tf.Tensor([5041  64], shape=(2,), dtype=int32)
64 Epoch 1/20
65 158/158 0s 1ms/step - loss: 0.0458
66 Epoch 2/20
67 158/158 0s 1ms/step - loss: 0.0435
68 Epoch 3/20
69 158/158 0s 834us/step - loss: 0.0395
70 Epoch 4/20
71 158/158 0s 790us/step - loss: 0.0364
72 Epoch 5/20
73 158/158 0s 919us/step - loss: 0.0383
74 Epoch 6/20
75 158/158 0s 743us/step - loss: 0.0367
76 Epoch 7/20
77 158/158 0s 846us/step - loss: 0.0382
78 Epoch 8/20
79 158/158 0s 788us/step - loss: 0.0362
80 Epoch 9/20
81 158/158 0s 757us/step - loss: 0.0359
82 Epoch 10/20
83 158/158 0s 850us/step - loss: 0.0355

```

```

91 | Epoch 11/20
92 | 158/158 0s 886us/step - loss: 0.0343
93 |
94 | Epoch 12/20
95 | 158/158 0s 752us/step - loss: 0.0345
96 |
97 | Epoch 13/20
98 | 158/158 0s 750us/step - loss: 0.0346
99 |
100 | Epoch 14/20
101 | 158/158 0s 841us/step - loss: 0.0334
102 |
103 | Epoch 15/20
104 | 158/158 0s 927us/step - loss: 0.0339
105 |
106 | Epoch 16/20
107 | 158/158 0s 813us/step - loss: 0.0275
108 |
109 | Epoch 17/20
110 | 158/158 0s 791us/step - loss: 0.0294
111 |
112 | Restoring model weights from the end of the best epoch: 19.
113 | model.metrics_names: ['loss']
114 | 40/40 0s 764us/step - loss: 0.0204
115 | [0.2460017055273056, 0.026489466428756714, 0.03945173695683479, 0.02035539597272873]
116 | processing fold # 4
117 | tf.Tensor([5041 64], shape=(2,), dtype=int32)
118 | Epoch 1/20
119 | 158/158 0s 872us/step - loss: 0.0234
120 | Epoch 2/20
121 | 158/158 0s 845us/step - loss: 0.0226
122 | Epoch 3/20
123 | 158/158 0s 978us/step - loss: 0.0246
124 | Epoch 4/20
125 | 158/158 0s 873us/step - loss: 0.0237
126 | Epoch 5/20
127 | 158/158 0s 803us/step - loss: 0.0226
128 | Epoch 6/20
129 | 158/158 0s 920us/step - loss: 0.0241
130 | Epoch 7/20
131 | 158/158 0s 927us/step - loss: 0.0247
132 | Epoch 8/20
133 | 158/158 0s 911us/step - loss: 0.0219
134 | Epoch 9/20
135 | 158/158 0s 843us/step - loss: 0.0219
136 | Epoch 10/20
137 | 158/158 0s 815us/step - loss: 0.0222
138 | Epoch 11/20
139 | 158/158 0s 917us/step - loss: 0.0216
140 | Epoch 12/20
141 | 158/158 0s 956us/step - loss: 0.0212
142 | Epoch 13/20
143 | 158/158 0s 933us/step - loss: 0.0230
144 | Epoch 14/20
145 | 158/158 0s 906us/step - loss: 0.0230
146 | Epoch 15/20
147 | 158/158 0s 843us/step - loss: 0.0223
148 | Epoch 16/20
149 | 158/158 0s 905us/step - loss: 0.0209
150 | Epoch 17/20
151 | 158/158 0s 860us/step - loss: 0.0212
152 | Epoch 18/20
153 | 158/158 0s 792us/step - loss: 0.0221
154 | Epoch 19/20
155 | 158/158 0s 896us/step - loss: 0.0237
156 | Epoch 20/20
157 | 158/158 0s 949us/step - loss: 0.0225
158 | Restoring model weights from the end of the best epoch: 16.
159 | model.metrics_names: ['loss']
160 | 40/40 0s 754us/step - loss: 0.1070
161 | [0.2460017055273056, 0.026489466428756714, 0.03945173695683479, 0.02035539597272873, 0.1070471853017807]
162 |
163 |
164 |
165 | Part 1 Kendall Tau (Gnutella 08): 0.8481805105261698
166 |
167 |
168 |
169 | WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model(model)'. This file
170 | format is considered legacy. We recommend using instead the native Keras format, e.g. 'model.save('my_model.keras')'
171 | or 'keras.saving.save_model(model, 'my_model.keras')'.
172 | Loading cached betweenness centrality for Gnutella 04...
173 | 340/340 0s 1ms/step
174 |
175 | Part 2 Kendall Tau (Gnutella 04): 0.6727189830593482

```

C Search Engine Code

Below is the implementation of the PageRank algorithm:

Listing 3: src/pagerank.py

```
"""
PageRank algorithm implementation for ranking web pages.

Recall from lecture that the PageRank algorithm is one popular example of an
algorithm that can be applied to rank web pages. When returning search results,
we want to rank pages that are more relevant to the query higher than pages
that are less relevant, and so we turn to algorithms like PageRank to help us
do this.

The PageRank algorithm assigns importance scores to pages based on the link
structure of the web graph. Pages that are linked to by many important pages
receive higher scores.

In this assignment, you will implement the PageRank algorithm yourself below
in the compute_pagerank function. This should be the ONLY implementation
necessary for this part of the assignment. You should NOT call a PageRank
implementation from external libraries (e.g. NetworkX, etc.).

However, PageRank is just one technique we can use to rank search results.
For part of the optional extra credit portion of this question, we invite
you to try out other techniques to rank search results (e.g. your own
ML-based approach, etc.).

"""

def compute_pagerank(graph, damping=0.85, max_iter=10000, tol=1e-8):
    """
    Compute PageRank scores for all nodes in the graph.

    Args:
        graph: Dictionary mapping URLs to lists of outbound links.
               Format: {url: [list of linked URLs]}
        damping: Damping factor (default 0.85). Probability of following
                 a link vs. jumping to a random page.
        max_iter: Maximum number of iterations (default 1000)
        tol: Convergence tolerance (default 1e-7)

    Returns:
        Dictionary mapping URLs to their PageRank scores.
        Format: {url: pagerank_score}
    """

    all_nodes = set(graph.keys())
    for links in graph.values():
        all_nodes.update(links)
    all_nodes = list(all_nodes)
    n = len(all_nodes)

    if n == 0:
        return {}

    # Initialize PageRank uniformly
    pr = {node: 1.0 / n for node in all_nodes}
```

```

54
55     for _ in range(max_iter):
56         new_pr = {node: 0.0 for node in all_nodes}
57         sink_pr_sum = 0.0
58
59         # Calculate contribution from sink nodes
60         for node in all_nodes:
61             out_links = graph.get(node, [])
62             if not out_links:
63                 sink_pr_sum += pr[node]
64
65         # Distribute mass from nodes with links
66         for node in all_nodes:
67             out_links = graph.get(node, [])
68             if out_links:
69                 share = (damping * pr[node]) / len(out_links)
70                 for target in out_links:
71                     if target in new_pr:
72                         new_pr[target] += share
73
74         # Add damping factor (teleportation) and sink distribution
75         # Total mass to distribute to each node from sinks and random jumps
76         base_add = ((1.0 - damping) + damping * sink_pr_sum) / n
77
78         diff = 0.0
79         for node in all_nodes:
80             new_pr[node] += base_add
81             diff += abs(new_pr[node] - pr[node])
82
83         pr = new_pr
84
85         if diff < tol:
86             break
87
88     return pr
89
90
91 if __name__ == "__main__":
92     # Simple test case
93     network1 = {
94         'A': ['B', 'C'],
95         'B': ['C'],
96         'C': ['A'],
97         'D': ['C']
98     }
99
100    # Slightly less trivial test case
101    network2 = {
102        'A': ['B', 'C', 'D'],
103        'B': ['E'],
104        'C': ['E'],
105        'D': ['E'],
106        'E': []
107    }
108
109    # Add more networks here if desired
110
111    test_cases = [
112        ("Network1", network1),

```

```

113     ("Network2", network2)
114 ]
115
116 print("\n" + "=" * 48)
117 print("Running PageRank on Simple Networks")
118 print("=" * 48 + "\n")
119
120 for name, g in test_cases:
121     print(f"{name}")
122     print("-" * len(name))
123     # Run PageRank on the network with default parameters
124     scores = compute_pagerank(g)
125     if scores:
126         print("{:<6}|{:>10}".format("Node", "PageRank"))
127         print("-" * 21)
128         for url, score in sorted(scores.items(), key=lambda x: x[1],
129                                   reverse=True):
130             print("{:<6}|{:>10.4f}".format(url, score))
131         print()
132     else:
133         print("PageRank not yet implemented!\n")

```